

Going low level with TCP sockets and :gen_tcp

Orestis Markou

@orestis

orestis.gr

https://github.com/orestis/elixir_tcp

Internet Protocol (IP in TCP/IP)

- Like passing notes in school
- Put data in a packet, pass it on
- Hope for the best!

Transport Control Protocol (TCP in TCP/IP)

- Point-to-point,
- stream
- two-way



BSD Sockets API

- `socket()`
- `bind()`
- `listen()`
- `accept()`
- `connect()`
- `gethostbyname()`
- ...

: gen_tcp

- **accept/1,2**
- **close/1**
- **connect/3,4**
- **listen/2**
- **recv/2,3**
- **send/2**
- **shutdown/2**
- **controlling_process/2**

:inet

- **:inet.gethostbyname/1,2**
- **:inet.setopts/2**
- ...

Hello world, server

- Accept connections on port 4001
- Send the current datetime
- Close the connection

Demo

example1.exs

Server Code

```
def server do
  {:ok, listen_socket} = :gen_tcp.listen(4001, [:binary,
                                                reuseaddr: true])

  server_handler(listen_socket)
end

def server_handler(listen_socket) do
  {:ok, socket} = :gen_tcp.accept(listen_socket)
  d = DateTime.utc_now() |> DateTime.to_string()
  :ok = :gen_tcp.send(socket, d <> "\r\n")
  :ok = :gen_tcp.shutdown(socket, :read_write)
  server_handler(listen_socket)
end
```

Server Code

```
def server do
  {:ok, listen_socket} = :gen_tcp.listen(4001, [:binary,
                                                reuseaddr: true])

  server_handler(listen_socket)
end

def server_handler(listen_socket) do
  {:ok, socket} = :gen_tcp.accept(listen_socket)
  d = DateTime.utc_now() |> DateTime.to_string()
  :ok = :gen_tcp.send(socket, d <> "\r\n")
  :ok = :gen_tcp.shutdown(socket, :read_write)
  server_handler(listen_socket)
end
```

Server Code

```
def server do
  {:ok, listen_socket} = :gen_tcp.listen(4001, [:binary,
                                                reuseaddr: true])

  server_handler(listen_socket)
end
```

```
def server_handler(listen_socket) do
  {:ok, socket} = :gen_tcp.accept(listen_socket)
  d = DateTime.utc_now() |> DateTime.to_string()
  :ok = :gen_tcp.send(socket, d <> "\r\n")
  :ok = :gen_tcp.shutdown(socket, :read_write)
  server_handler(listen_socket)
end
```

Client code

```
def client do
  {:ok, socket} = :gen_tcp.connect('localhost', 4001,
                                   [:binary, active: true])
  client_handler(socket)
end

def client_handler(socket) do
  receive do
    {:tcp, ^socket, data} ->
      IO.write data
      client_handler(socket)
    {:tcp_closed, ^socket} -> IO.puts "== CLOSED =="
  end
end
```

Client code

```
def client do
  {:ok, socket} = :gen_tcp.connect('localhost', 4001,
                                   [:binary, active: true])
  client_handler(socket)
end

def client_handler(socket) do
  receive do
    {:tcp, ^socket, data} ->
      IO.write data
      client_handler(socket)
    {:tcp_closed, ^socket} -> IO.puts "== CLOSED =="
  end
end
```

Client code

```
def client do
  {:ok, socket} = :gen_tcp.connect('localhost', 4001,
                                   [:binary, active: true])
  client_handler(socket)
end

def client_handler(socket) do
  receive do
    {:tcp, ^socket, data} ->
      IO.write data
      client_handler(socket)
    {:tcp_closed, ^socket} -> IO.puts "== CLOSED =="
  end
end
```

One server, many clients?

Server Code

```
def server do
  {:ok, listen_socket} = :gen_tcp.listen(4001, [:binary,
                                                reuseaddr: true])

  for _ <- 0..10, do: spawn(fn -> server_handler(listen_socket) end)
  Process.sleep(:infinity)
end

def server_handler(listen_socket) do
  {:ok, socket} = :gen_tcp.accept(listen_socket)
  :ok = :gen_tcp.send(socket, "Hello!\r\n")
  :ok = :gen_tcp.shutdown(socket, :read_write)
  server_handler(listen_socket)
end
```


Two-way stream - server

- Accept connections
- Send "HELLO?"
- Wait for name
- Send "Hello, <name>!"
- Close the connection

Two-way stream - client

- Connect
- Wait for "**HELLO?**"
- Send name
- Read all data
- Wait until connection is closed

Two way - demo

example2b.exs

Two-way - server

```
def server_handler(listen_socket) do
  {:ok, socket} = :gen_tcp.accept(listen_socket)
  :ok = :gen_tcp.send(socket, "HELLO?")
  receive do
    {:tcp, ^socket, data} ->
      :ok = :gen_tcp.send(socket, "Hello, #{data}!\r\n")
  end
  :ok = :gen_tcp.shutdown(socket, :read_write)
  server_handler(listen_socket)
end
```

Two-way - client

```
def client_handler(socket) do
  receive do
    {:tcp, ^socket, "HELLO?"} ->
      d = IO.gets("Enter your name: ") |> String.trim()
      :ok = :gen_tcp.send(socket, d)
      client_handler(socket)
    {:tcp, ^socket, data} ->
      IO.write data
      client_handler(socket)
    {:tcp_closed, ^socket} -> IO.puts "== CLOSED =="
  end
end
```

Two-way - client

```
def client_handler(socket) do
  receive do
    {:tcp, ^socket, "HELLO?"} ->
      d = IO.gets("Enter your name: ") |> String.trim()
      :ok = :gen_tcp.send(socket, d)
      client_handler(socket)
    {:tcp, ^socket, data} ->
      IO.write data
      client_handler(socket)
    {:tcp_closed, ^socket} -> IO.puts "== CLOSED =="
  end
end
```

Passive mode

- Closer to the original BSD API
- Read/write to a "file"
- Blocking API with timeouts
- Provides back-pressure

Passive mode server

```
def server do
  {:ok, listen_socket} = :gen_tcp.listen(4001, [:binary, reuseaddr: true
                                              active: false])

  server_handler(listen_socket)
end
```

```
def server_handler(listen_socket) do
  {:ok, socket} = :gen_tcp.accept(listen_socket)
  :ok = :gen_tcp.send(socket, "HELLO?")
  {:ok, data} = :gen_tcp.recv(socket, 0, 5000)
  :ok = :gen_tcp.send(socket, "Hello, #{data}!\r\n")
  :ok = :gen_tcp.shutdown(socket, :read_write)
  server_handler(listen_socket)
end
```


Passive mode client

```
def client do
  {:ok, socket} = :gen_tcp.connect('localhost', 4001,
    [:binary,
     active: false])
  client_handler(socket)
end
```

Passive mode client

```
def client_handler(socket) do
  case :gen_tcp.recv(socket, 0, 5000) do
    {:ok, "HELLO?"} ->
      d = IO.gets("Enter your name: ") |> String.trim()
      :ok = :gen_tcp.send(socket, d)
      client_handler(socket)
    {:ok, data} ->
      IO.write data
      client_handler(socket)
    {:error, :closed} -> IO.puts "== CLOSED =="
  end
end
```

Passive mode

- `recv(socket, length)`
- `recv(socket, length, timeout)`
- `timeout` defaults to `:infinity`
- When `length > 0`, "read exactly `length` bytes"
- When `length == 0`, "read all available"

Hello HTTP (whoops)

```
{:ok, socket} = :gen_tcp.connect('www.gutenberg.org', 80,  
                                  [:binary, active: false])  
  
:ok = :gen_tcp.send(socket,  
                    ["GET /files/84/84-0.txt HTTP/1.0\r\n",  
                    "Host: www.gutenberg.org\r\n",  
                    "Accept: text/plain\r\n\r\n"])  
  
{:ok, response} = :gen_tcp.recv(socket, 0, 5000)  
IO.puts response  
IO.puts "==== Received #{byte_size(response)} bytes ===="
```

Major gotchas

- *Passive:* How many bytes to read?
- *Active:* Will "HELLO?" arrive in a single message?
- **This is by design!**

Protocols

- Give shape to the data packets
- Common or niche or your own!

Protocol specifications

- What comes next?
- What form does it come in?
- Who is responsible for the next transmission?
- (Distributed state machine)

Protocol specifications

e.g. Daytime protocol (RFC 867)

TCP Based Daytime Service

One daytime service is defined as a connection based application on TCP. A server listens for TCP connections on TCP port 13. Once a connection is established the current date and time is sent out the connection as a ascii character string (and any data received is thrown away). The service closes the connection after sending the quote.

Protocol specifications

e.g. HTTP/1.1 protocol (RFC 2616)

<176 pages>

e.g. Memcached protocol

<1200 lines>

Hello HTTP *

```
{:ok, socket} = :gen_tcp.connect('www.gutenberg.org', 80,  
                                  [:binary, active: false])  
  
:ok = :gen_tcp.send(socket,  
                    ["GET /files/84/84-0.txt HTTP/1.0\r\n",  
                    "Host: www.gutenberg.org\r\n",  
                    "Accept: text/plain\r\n\r\n"])  
  
response = _recv(socket, [])  
IO.puts response  
IO.puts "==== Received #{byte_size(response)} bytes ===="
```

```
def _recv(socket, acc) do
  r = :gen_tcp.recv(socket, 0, 5000)
  case r do
    {:ok, data} -> _recv(socket, [data|acc])
    other ->
      IO.puts("other")
      IO.inspect(other)
      Enum.reverse(acc) |> IO.iodata_to_binary()
  end
end
```

Built-in protocols

- Provided by :gen_tcp
- Limited in scope, non-extensible
- Might be useful

Prefix header length

[packet: 2]

- Transparently add/strip header
- 1, 2 or 4 byte header length
- Support up to 2GB messages
- Very useful when you control both ends
- Use **0** for "raw" mode (default)

Line-based messages

```
[packet: :line,  
line_delimiter: ?\n,  
packet_size: 255]
```

- Split incoming messages by newline
- Outgoing messages are your responsibility
- A few gotchas, must evaluate

Mutable sockets

- Can change mode on-the-fly (binary, active)
- Active mode can be one shot or N-shot or permanent
- Can change protocols on the fly
- Read a line, extract content length, read raw bytes

Demo

protocols.exs

/usr/local/opt/memcached/bin/memcached


```
def memcached_client_get do
  {:ok, socket} = :gen_tcp.connect('localhost', 11211,
    [:binary, active: false,
     packet: :line])
  :gen_tcp.send(socket, "get elixirconf\r\n")
  {:ok, response} = :gen_tcp.recv(socket, 0, 5000)
  IO.puts "Raw response:"
  IO.inspect response
  <<"VALUE elixirconf ", resp::binary>> = response
  [_ , length] = resp |> String.trim() |> String.split()
    |> Enum.map(&String.to_integer/1)
  :inet.setopts(socket, [packet: 0])
  {:ok, data} = :gen_tcp.recv(socket, length, 5000)
  IO.puts "Actual data:"
  IO.inspect data
end
```

```
def memcached_client_get do
  {:ok, socket} = :gen_tcp.connect('localhost', 11211,
    [:binary, active: false,
      packet: :line])
  :gen_tcp.send(socket, "get elixirconf\r\n")
  {:ok, response} = :gen_tcp.recv(socket, 0, 5000)
  IO.puts "Raw response:"
  IO.inspect response
  <<"VALUE elixirconf ", resp::binary>> = response
  [_ , length] = resp |> String.trim() |> String.split()
    |> Enum.map(&String.to_integer/1)
  :inet.setopts(socket, [packet: 0])
  {:ok, data} = :gen_tcp.recv(socket, length, 5000)
  IO.puts "Actual data:"
  IO.inspect data
end
```

Pain point: Untangle the protocol logic from the socket logic

- Abstract the "transport" out
- Provide a dummy transport for testing
- Transparently adapt to TLS/SSL, tunnels etc.
- Timeout handling

Thank you!

- https://github.com/orestis/elixir_tcp
- TCP/IP Illustrated, Volume 1 [Fall & Stevens]
- http://erlang.org/doc/man/gen_tcp.html
- <http://erlang.org/doc/man/inet.html>
- <https://ninenines.eu/docs/en/ranch/1.4/guide/>

Orestis Markou

@orestis

orestis.gr

